


Compute cost on 100K reads (single H200)

End-to-end speed



Inference memory

